

# Sales Forecasting and Inventory Insights with Machine Learning: Synthetic Data Case Study

*A practical demonstration of demand forecasting using time series models and strategic inventory risk analysis with synthetic retail data.*

**Author:** Jose Araújo

**Date:** 25/09/2025

**GitHub:** <https://github.com/josefaquino>

## Technical Development Plan for ML-Based Demand Forecasting and Inventory Optimization Project

### Month 1: Foundations of Forecasting and Monitoring

#### 1. Demand Forecasting Model

- Develop time series models that account for seasonality, stockouts, and competitor behavior.
- Recommended techniques:
  - Hybrid models: Prophet combined with tree-based methods like LightGBM to incorporate external variables.
  - Deep learning models (LSTM, Transformers with attention) based on data volume and infrastructure.
  - Advanced feature engineering including holidays, promotions, market events, and competitor data.

#### 2. Safety Stock Calculator

- Integrate forecast uncertainty and lead-time variability to calculate adaptive safety stock levels.
- Use probabilistic models and Monte Carlo simulations to assess risks and define inventory buffers.

#### 3. Daily Stockout Risk Alert System

- Build an automated system that monitors forecast and inventory status to generate early stockout alerts.
- Implement APIs or dashboards presenting key risk indicators to the purchasing team.

### Month 2: Competitiveness and Market Monitoring

#### 1. Buy Box Share Prediction Model

- Create predictive models to estimate Buy Box winning probability on Amazon based on price, availability, and competitor signals.
- Apply classification and scoring techniques like XGBoost or neural networks with categorical embeddings.

#### 2. Competitor Inventory Monitoring

- Develop data pipelines for frequent collection of competitor inventory via scraping or APIs.
- Incorporate competitor inventory data into forecasting and alert systems.

#### 3. Strategic Purchase Opportunity Scanner

- Analyze market gaps and competitor stock levels to recommend strategic purchase opportunities with higher sales potential.

### Month 3: Advanced Optimization and Automation

#### 1. Price Elasticity Models by Product Category

- Analyze how price changes affect demand elasticity per category, using regression and causal inference models.

#### 2. Automated Purchasing Decision Engine

- Implement a decision system using business rules and reinforcement learning to optimize replenishment balancing cost, risk, and demand.

#### 3. Full Integration

- Seamlessly integrate forecasting, optimization, and alert tools with existing company systems and provide thorough user documentation.

This structured plan aims to deliver impactful, scalable, and production-ready ML solutions tailored to the complexities of a large-scale Amazon retail inventory environment.

```
# -*- coding: utf-8 -*-
# Colab code for Forecasting and Insights with Synthetic Data

import pandas as pd
import numpy as np
from datetime import datetime, timedelta
from prophet import Prophet
import matplotlib.pyplot as plt
from google.colab import files

# --- Generate synthetic data ---

num_skus = 20
periods = 180 # 6 months of data
start_date = datetime.today() - timedelta(days=periods)
dates = pd.date_range(start=start_date, periods=periods)

np.random.seed(42)
rows = []
for sku in range(1, num_skus + 1):
    base_demand = np.random.randint(50, 200) # average daily demand
    seasonality = 10 * np.sin(np.linspace(0, 3 * np.pi, periods)) # simple seasonality
    noise = np.random.normal(0, 10, periods) # noise
    demand = base_demand + seasonality + noise
    demand = np.clip(demand, 0, None) # demand is not negative

    stockouts = np.random.choice([0, 1], size=periods, p=[0.05, 0.95]) # 5% stockouts
    demand = demand * stockouts

    for date, sales in zip(dates, demand):
        rows.append({'date': date, 'sku': f'SKU_{sku}', 'sales': sales})

synthetic_data = pd.DataFrame(rows)

# Save CSV file for download
synthetic_data.to_csv('synthetic_sales_data.csv', index=False)
files.download('synthetic_sales_data.csv')

# --- Forecast model for 1 SKU ---

# Use the first SKU for forecasting example
sku_id = synthetic_data['sku'].unique()[0]
sku_data = synthetic_data[synthetic_data['sku'] == sku_id][['date', 'sales']].rename(columns={'date': 'ds', 'sales': 'y'})

# Create and fit Prophet model
model = Prophet(yearly_seasonality=True, weekly_seasonality=True, daily_seasonality=False)
model.fit(sku_data)

# Make forecast for next 30 days
future = model.make_future_dataframe(periods=30)
forecast = model.predict(future)

# Plot forecast
fig1 = model.plot(forecast)
plt.title(f'Sales Forecast - {sku_id}')
plt.show()

fig2 = model.plot_components(forecast)
plt.show()

# --- Strategic insights ---

# 1) SKUs with most stockouts (zero sales days)
#ruptures = synthetic_data.groupby('sku').apply(lambda x: (x['sales'] == 0).sum()).sort_values(ascending=False)
ruptures = synthetic_data.groupby('sku').apply(
    lambda x: (x['sales'] == 0).sum(),
    include_groups=False
).sort_values(ascending=False)

# 2) SKUs with highest demand variability (candidates for higher safety stock)
variability = synthetic_data.groupby('sku')['sales'].std().sort_values(ascending=False)
```

```
print("Top SKUs with risk of stockout (most zero sales days):")
print(ruptures.head(5))

print("\nTop SKUs with highest demand variability:")
print(variability.head(5))
```

```

import pandas as pd
import numpy as np
from prophet import Prophet
from sklearn.metrics import mean_absolute_error, mean_squared_error
import matplotlib.pyplot as plt

# Load synthetic data previously generated
df = synthetic_data.copy()

# Prepare data for Prophet: choose one SKU as example
sku_id = df['sku'].unique()[0]
sku_data = df[df['sku'] == sku_id][['date', 'sales']].rename(columns={'date': 'ds', 'sales': 'y'})

# Add external regressors simulating competitor price index (synthetic)
np.random.seed(42)
# Generate competitor price index for historical period
comp_price_index_hist = 1 + 0.1 * np.sin(np.linspace(0, 3 * np.pi, len(sku_data))) + np.random.normal(0, 0.02, len(sku_data))
sku_data['comp_price_index'] = comp_price_index_hist

# Initialize Prophet with extra regressors
model = Prophet(yearly_seasonality=True, weekly_seasonality=True)
model.add_regressor('comp_price_index')

# Fit model
model.fit(sku_data)

# Create future dataframe including 30 days for prediction
future = model.make_future_dataframe(periods=30)

# Generate competitor price index for full period (historical + future)
full_comp_price_index = 1 + 0.1 * np.sin(np.linspace(0, 3 * np.pi + np.pi/6, len(future))) + np.random.normal(0, 0.02, len(future))
full_comp_price_df = pd.DataFrame({'ds': future['ds'], 'comp_price_index': full_comp_price_index})

# Merge competitor price index into future dataframe, ensure no NaNs
future = future.drop(columns=['comp_price_index'], errors='ignore')
future = pd.merge(future, full_comp_price_df, on='ds', how='left')
future['comp_price_index'] = future['comp_price_index'].ffill()




# Predict
forecast = model.predict(future)


```

```
# Evaluate on known period (hold-out last 30 days)
eval_true = sku_data['y'].tail(30).values
eval_pred = forecast['yhat'].head(len(sku_data)).tail(30).values
mae = mean_absolute_error(eval_true, eval_pred)
rmse = np.sqrt(mean_squared_error(eval_true, eval_pred))
print(f'Mean Absolute Error (MAE): {mae:.2f}')
print(f'Root Mean Squared Error (RMSE): {rmse:.2f}')

# Plot forecast and components
fig1 = model.plot(forecast)
plt.title(f'Demand Forecast with Competitor Price Index - {sku_id}')
plt.show()

fig2 = model.plot_components(forecast)
plt.show()
```





Top SKUs with risk of stockout (most zero sales days):

```
sku
SKU_3    18
SKU_11   15
SKU_7    14
SKU_6    13
SKU_1    12
dtype: int64
```

Top SKUs with highest demand variability:


```
sku
SKU_11   50.074841
SKU_10   47.956108
SKU_14   40.539771
SKU_1    40.360021
SKU_4    40.136372
Name: sales, dtype: float64
```

```

INFO:prophet:Disabling daily seasonality. Run prophet with daily_seasonality=True to override this.
DEBUG:cmdstanpy:input tempfile: /tmp/tmp890y_d3f/cwzca00z.json
DEBUG:cmdstanpy:input tempfile: /tmp/tmp890y_d3f/vtg8qvz8.json
DEBUG:cmdstanpy:idx 0
DEBUG:cmdstanpy:running CmdStan, num_threads: None
DEBUG:cmdstanpy:CmdStan args: ['/usr/local/lib/python3.12/dist-packages/prophet/stan_model/prophet_model.bin', 'random
02:03:38 - cmdstanpy - INFO - Chain [1] start processing
INFO:cmdstanpy:Chain [1] start processing
02:03:38 - cmdstanpy - INFO - Chain [1] done processing
INFO:cmdstanpy:Chain [1] done processing
Mean Absolute Error (MAE): 26.68
Root Mean Squared Error (RMSE): 47.56

```

Demand Forecast with Competitor Price Index - SKU\_1



```
import numpy as np
```

```

# Parameters for Safety Stock Calculation
service_level = 0.95 # Desired service level (95%)
Z = 1.65 # z-score for 95% service level (from standard normal distribution)

# Assume lead time parameters (in days)
lead_time_mean = 7 # average lead time of 7 days
lead_time_std = 2 # std deviation of lead time is 2 days

# Calculate demand statistics from historical data
# Using the same SKU data as before
demand_mean = sku_data['y'].mean() # average daily demand
demand_std = sku_data['y'].std() # standard deviation daily demand

# Safety stock formula considering demand and lead time variability
safety_stock = Z * np.sqrt((demand_std ** 2) * lead_time_mean + (demand_mean ** 2) * (lead_time_std ** 2))

print(f"Average daily demand: {demand_mean:.2f}")
print(f"Demand standard deviation: {demand_std:.2f}")
print(f"Lead time mean (days): {lead_time_mean}")
print(f"Lead time std deviation (days): {lead_time_std}")
print(f"Calculated safety stock (units): {safety_stock:.2f}")

```

```

Average daily demand: 143.65
Demand standard deviation: 26.68
Lead time mean (days): 7
Lead time std deviation (days): 2
Calculated safety stock (units): 505.72

```

```

# Assume we have the current stock level for the SKU under analysis (or for all SKUs)
current_stock = 450 # Example: current stock in units


# Compare with the calculated safety stock
if current_stock < safety_stock:
    print(f"ALERT: Stock for {sku_id} is below safety stock level! Current stock: {current_stock:.0f}, Safety stock: {safety_stock:.0f}")
else:
    print(f"Stock level for {sku_id} is healthy. Current stock: {current_stock:.0f}, Safety stock: {safety_stock:.0f}")

```

```

ALERT: Stock for SKU_1 is below safety stock level! Current stock: 450, Safety stock: 506
Sunday Monday Tuesday Wednesday Thursday Friday Saturday
Day of week

```



This simple script can be adapted to run daily with real inventory data to generate alerts via email, dashboards, or reporting tools.



```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from xgboost import XGBClassifier
from sklearn.metrics import roc_auc_score, classification_report

# Generate synthetic dataset for Buy Box prediction
np.random.seed(42)
num_samples = 1000

data = pd.DataFrame({
    'price_diff': np.random.normal(0, 10, num_samples), # price difference to competitor (negative means cheaper)
    'inventory_level': np.random.randint(0, 500, num_samples), # current inventory
    'fba': np.random.choice([0, 1], num_samples), # Fulfilled by Amazon (1) or not (0)
    'seller_rating': np.random.uniform(3.5, 5.0, num_samples), # seller rating (3.5 to 5)
    'promotion': np.random.choice([0, 1], num_samples), # promotion running (1) or not (0)
    'buy_box_win': np.random.choice([0, 1], num_samples, p=[0.7, 0.3]) # target variable, 1 means won buy box
})

# Features and target
X = data.drop('buy_box_win', axis=1)
y = data['buy_box_win']

# Train/Test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y, random_state=42)

# Feature scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# XGBoost Classifier
model = XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_state=42)
model.fit(X_train_scaled, y_train)

# Predict probabilities
y_pred_proba = model.predict_proba(X_test_scaled)[ :, 1]
y_pred = model.predict(X_test_scaled)

# Evaluation
auc = roc_auc_score(y_test, y_pred_proba)
print(f'ROC AUC Score: {auc:.3f}\n')

print("Classification Report:")
print(classification_report(y_test, y_pred))

# Feature importance plot (optional)
import matplotlib.pyplot as plt
import xgboost as xgb

xgb.plot_importance(model)
plt.title('Feature Importance - Buy Box Prediction')
plt.show()
```




```
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:183: UserWarning: [02:08:45] WARNING: /workspace/src/learn
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
ROC AUC Score: 0.425
```

Classification Report:

	precision	recall	f1-score	support
0	0.72	0.90	0.80	146
1	0.21	0.07	0.11	54
accuracy			0.68	200
macro avg	0.47	0.49	0.46	200
weighted avg	0.59	0.68	0.61	200

Feature Importance - Buy Box Prediction



```
import pandas as pd
import numpy as np

np.random.seed(42)

# Parameters
num_days = 2 * 365 + 30 # approximately 2 years and 1 month
dates = pd.date_range(start='2023-01-01', periods=num_days)

# Simulation output dataframe
data = pd.DataFrame({'date': dates})

# Annual seasonality (sine) + weekly trends + noise
data['base_sales'] = 50 + 20 * np.sin(2 * np.pi * data.index / 365) + 5 * np.sin(2 * np.pi * data.index / 7)
data['base_sales'] += np.random.normal(0, 5, num_days)
data['base_sales'] = np.clip(data['base_sales'], 0, None)

# Competitor price fluctuates around 100, with noise and trend
data['comp_price'] = 100 + 5 * np.sin(2 * np.pi * data.index / 180) + np.random.normal(0, 2, num_days)

# Our pricing strategy (price relative to competitor)
data['price_diff'] = np.random.normal(-1, 2, num_days) # price slightly lower than competitor

# Daily stock levels simulating declines and replenishments
stock = 300
stock_levels = []
for i in range(num_days):
    demand = max(0, int(data.loc[i, 'base_sales'] + np.random.normal(0, 3)))
    stock = max(0, stock - demand)
    if i % 14 == 0: # replenishment every 2 weeks
        stock += 200 + np.random.randint(-20, 20)
    stock_levels.append(stock)
data['stock_level'] = stock_levels

# Random promotions on 15% of days
data['promotion'] = (np.random.rand(num_days) < 0.15).astype(int)

# Seller rating with small variation
data['seller_rating'] = np.clip(np.random.normal(4.7, 0.1, num_days), 3.5, 5.0)

# Basic probability of winning the Buy Box based on price and stock
data['buy_box_prob'] = (
    0.3 * (data['price_diff'] < 0).astype(float) +
    0.4 * (data['stock_level'] > 50).astype(float) +
    0.1 * data['promotion'] +
    0.2 * ((data['seller_rating'] - 3.5) / 1.5)
)
data['buy_box_prob'] = data['buy_box_prob'].clip(0, 1)

# Simulate whether Buy Box was won on that day
data['buy_box_win'] = (np.random.rand(num_days) < data['buy_box_prob']).astype(int)

# Select features for modeling
```

```
data_model = data[[
    'date', 'price_diff', 'stock_level', 'promotion', 'seller_rating', 'buy_box_win'
]]

print(data_model.head())

# This dataset can now be used to train the Buy Box Share prediction model
```

	date	price_diff	stock_level	promotion	seller_rating	buy_box_win
0	2023-01-01	1.442067	469	0	4.679627	0
1	2023-01-02	0.164195	418	0	4.880652	0
2	2023-01-03	-1.452968	357	0	4.664929	1
3	2023-01-04	-2.918878	295	0	4.804999	1
4	2023-01-05	-1.744414	249	0	4.733889	1

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from xgboost import XGBClassifier
from sklearn.metrics import roc_auc_score, classification_report
import matplotlib.pyplot as plt
import xgboost as xgb

# Assume data_model is your synthetic dataset from before
data = data_model.copy()

# Features and target
X = data.drop(['date', 'buy_box_win'], axis=1)
y = data['buy_box_win']

# Train/Test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y, random_state=42)

# Feature scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# XGBoost Classifier training
model = XGBClassifier(eval_metric='logloss', random_state=42)
model.fit(X_train_scaled, y_train)

# Predict probabilities and classes
y_pred_proba = model.predict_proba(X_test_scaled)[: , 1]
y_pred = model.predict(X_test_scaled)

# Evaluate
auc = roc_auc_score(y_test, y_pred_proba)
print(f'ROC AUC Score: {auc:.3f}\n')

print("Classification Report:")
print(classification_report(y_test, y_pred))


# Feature importance plot
xgb.plot_importance(model)
plt.title('Feature Importance - Buy Box Prediction')
plt.show()
```

ROC AUC Score: 0.665

Classification Report:

	precision	recall	f1-score	support
0	0.67	0.57	0.62	77
1	0.62	0.71	0.66	75
accuracy			0.64	152
macro avg	0.64	0.64	0.64	152
weighted avg	0.64	0.64	0.64	152

Feature Importance - Buy Box Prediction



```
from imblearn.over_sampling import SMOTE
from sklearn.pipeline import Pipeline

# Balancear dados usando SMOTE
smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X, y)

# Dividir dados balanceados em treino e teste
X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled,
                                                    test_size=0.2, random_state=42, stratify=y_resampled)

# Escalar características
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Treinar modelo XGBoost
model = XGBClassifier(eval_metric='logloss', random_state=42)
model.fit(X_train_scaled, y_train)

# Avaliar
y_pred_proba = model.predict_proba(X_test_scaled)[:,-1]
y_pred = model.predict(X_test_scaled)

auc = roc_auc_score(y_test, y_pred_proba)
print(f'ROC AUC Score after SMOTE: {auc:.3f}\n')
print("Classification Report after SMOTE:")
print(classification_report(y_test, y_pred))
```

ROC AUC Score after SMOTE: 0.674

Classification Report after SMOTE:

	precision	recall	f1-score	support
0	0.61	0.56	0.58	77
1	0.59	0.64	0.61	77
accuracy			0.60	154
macro avg	0.60	0.60	0.60	154
weighted avg	0.60	0.60	0.60	154

```
from sklearn.model_selection import RandomizedSearchCV
from xgboost import XGBClassifier
import numpy as np

# Parâmetros para tunagem
param_dist = {
    'n_estimators': [50, 100, 200, 300],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.01, 0.05, 0.1, 0.2],
    'subsample': [0.6, 0.8, 1.0],
    'colsample_bytree': [0.6, 0.8, 1.0],
    'reg_alpha': [0, 0.1, 0.5],
    'reg_lambda': [1, 1.5, 2]
}
```

```

# Usar dados balanceados com SMOTE
smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X, y)

# Dividir em treino e teste
X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled, test_size=0.2, random_state=42, stratify=

# Escalar
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Modelo base
xgb_clf = XGBClassifier(eval_metric='logloss', use_label_encoder=False, random_state=42)

# Configurar busca randomizada - reduz tempo comparado ao grid search
random_search = RandomizedSearchCV(
    estimator=xgb_clf,
    param_distributions=param_dist,
    n_iter=30,
    scoring='roc_auc',
    cv=3,
    verbose=1,
    random_state=42,
    n_jobs=-1
)

# Executar busca
random_search.fit(X_train_scaled, y_train)

# Melhor modelo e parâmetros
best_model = random_search.best_estimator_
print("Best parameters:", random_search.best_params_)

# Avaliar no teste
y_pred_proba = best_model.predict_proba(X_test_scaled)[:,-1]
y_pred = best_model.predict(X_test_scaled)

auc = roc_auc_score(y_test, y_pred_proba)
print(f'ROC AUC Score after tuning: {auc:.3f}')
print(classification_report(y_test, y_pred))

```

Fitting 3 folds for each of 30 candidates, totalling 90 fits  
 Best parameters: {'subsample': 0.6, 'reg\_lambda': 1, 'reg\_alpha': 0.5, 'n\_estimators': 100, 'max\_depth': 7, 'learning\_rate': 0.1}  
 ROC AUC Score after tuning: 0.707

	precision	recall	f1-score	support
0	0.68	0.62	0.65	77
1	0.65	0.70	0.68	77
accuracy			0.66	154
macro avg	0.66	0.66	0.66	154
weighted avg	0.66	0.66	0.66	154

/usr/local/lib/python3.12/dist-packages/xgboost/training.py:183: UserWarning: [02:14:59] WARNING: /workspace/src/learn...  
 Parameters: { "use\_label\_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)
```

```

import pandas as pd
import numpy as np

np.random.seed(42)

# Simular dados de estoque de concorrentes para 100 SKUs ao longo de 60 dias
num_skus = 100
num_days = 60
dates = pd.date_range(start='2025-07-01', periods=num_days)

# Gerar DataFrame com estoque diário por SKU concorrente
data_list = []
for sku in range(1, num_skus + 1):
    stock_level = 500 + np.random.randint(-50, 50) # estoque inicial variado
    for day in dates:
        # Simula consumo e reabastecimento
        demand = max(0, int(np.random.normal(10, 3)))

```

```

stock_level = max(0, stock_level - demand)
if np.random.rand() < 0.1: # reposiç o ocasional
    stock_level += np.random.randint(50, 150)
data_list.append({'date': day, 'sku': f'SKU_{sku}', 'competitor_stock': stock_level})

competitor_stock_df = pd.DataFrame(data_list)

print(competitor_stock_df.head(10))

```

```

      date  sku  competitor_stock
0 2025-07-01  SKU_1             495
1 2025-07-02  SKU_1             622
2 2025-07-03  SKU_1             615
3 2025-07-04  SKU_1             744
4 2025-07-05  SKU_1             736
5 2025-07-06  SKU_1             725
6 2025-07-07  SKU_1             715
7 2025-07-08  SKU_1             711
8 2025-07-09  SKU_1             705
9 2025-07-10  SKU_1             695

```

```

import pandas as pd
import numpy as np
import plotly.express as px

np.random.seed(42)

# ----- Step 1: Generate synthetic competitor stock data -----
num_skus = 100
num_days = 60
dates = pd.date_range(start='2025-07-01', periods=num_days)

data_list = []
for sku in range(1, num_skus + 1):
    stock_level = 500 + np.random.randint(-50, 50) # initial stock varied
    for day in dates:
        demand = max(0, int(np.random.normal(10, 3))) # daily demand with noise
        stock_level = max(0, stock_level - demand)
        if np.random.rand() < 0.1: # occasional restocking
            stock_level += np.random.randint(50, 150)
        data_list.append({'date': day, 'sku': f'SKU_{sku}', 'competitor_stock': stock_level})

competitor_stock_df = pd.DataFrame(data_list)

# ----- Step 2: Calculate rolling average & std dev, define low stock alert -----
competitor_stock_df['stock_7d_avg'] = competitor_stock_df.groupby('sku')['competitor_stock'].transform(lambda x: x.rolling(7).mean())
competitor_stock_df['stock_7d_std'] = competitor_stock_df.groupby('sku')['competitor_stock'].transform(lambda x: x.rolling(7).std())
competitor_stock_df['low_stock_alert'] = competitor_stock_df['competitor_stock'] < (competitor_stock_df['stock_7d_avg'] - 2 * competitor_stock_df['stock_7d_std'])

# ----- Step 3: Generate synthetic own stock data -----
own_stock_data = []
for sku in range(1, num_skus + 1):
    stock_level = 300 + np.random.randint(-30, 30)
    for day in dates:
        stock_level = max(0, stock_level + np.random.randint(-10, 10)) # small fluctuations
        own_stock_data.append({'date': day, 'sku': f'SKU_{sku}', 'own_stock': stock_level})

df_own_stock = pd.DataFrame(own_stock_data)

# ----- Step 4: Merge dataframes for opportunity scanning -----
stock_comp = pd.merge(competitor_stock_df, df_own_stock, on=['date', 'sku'])

# Filter opportunities: competitor low stock AND own stock above threshold (e.g., 100)
opportunities = stock_comp[(stock_comp['low_stock_alert']) & (stock_comp['own_stock'] > 100)]

print("Sample opportunities:")
print(opportunities[['date', 'sku', 'competitor_stock', 'own_stock', 'low_stock_alert']].head(10))

# ----- Step 5: Visualization -----
# Aggregate average stock over time
daily_comp_stock = competitor_stock_df.groupby('date')['competitor_stock'].mean().reset_index()
daily_own_stock = df_own_stock.groupby('date')['own_stock'].mean().reset_index()
stock_comparison = pd.merge(daily_comp_stock, daily_own_stock, on='date', how='inner')

# Plot average stock levels over time
fig = px.line(stock_comparison, x='date', y=['competitor_stock', 'own_stock'],

```

```
labels={'value': 'Average Stock Level', 'date': 'Date', 'variable': 'Inventory'},
title='Average Competitor vs Own Stock Over Time')

fig.show()

# Plot competitor stock and low stock alerts for a specific SKU
sku_select = 'SKU_1'
sku_data = competitor_stock_df[competitor_stock_df['sku'] == sku_select].copy()


fig2 = px.line(sku_data, x='date', y='competitor_stock', title=f'Competitor Stock Over Time - {sku_select}')
fig2.add_scatter(x=sku_data[sku_data['low_stock_alert']]['date'],
                 y=sku_data[sku_data['low_stock_alert']]['competitor_stock'],
                 mode='markers', name='Low Stock Alert', marker=dict(color='red', size=10))

fig2.show()
```


sample opportunities:

	date	sku	competitor_stock	own_stock	low_stock_alert
9	2025-07-10	SKU_1	695	280	True
10	2025-07-11	SKU_1	681	286	True
11	2025-07-12	SKU_1	672	290	True
12	2025-07-13	SKU_1	664	284	True
13	2025-07-14	SKU_1	654	286	True
14	2025-07-15	SKU_1	646	279	True
15	2025-07-16	SKU_1	637	285	True
22	2025-07-23	SKU_1	701	293	True
23	2025-07-24	SKU_1	692	298	True
24	2025-07-25	SKU_1	680	303	True

Average Competitor vs Own Stock Over Time



Competitor Stock Over Time - SKU\_1



```
import plotly.express as px
import pandas as pd
```

```
# Assume competitor_stock_df and df_own_stock were created previously
```

```
# Aggregate stocks by date for overview
```

```
daily_competitor_stock = competitor_stock_df.groupby('date')['competitor_stock'].mean().reset_index()
daily_own_stock = df_own_stock.groupby('date')['own_stock'].mean().reset_index()
```

```



# Merge for comparative plotting
stock_comparison = pd.merge(daily_competitor_stock, daily_own_stock, on='date', how='inner')

# Plot average competitor and own stock levels over time
fig = px.line(stock_comparison, x='date', y=['competitor_stock', 'own_stock'],
              labels={'value': 'Average Stock Level', 'date': 'Date', 'variable': 'Inventory'},
              title='Average Competitor vs Own Stock Over Time')
fig.show()

# Visualize low stock alerts for a specific SKU example
sku_select = 'SKU_1'
sku_data = competitor_stock_df[competitor_stock_df['sku'] == sku_select].copy()

fig2 = px.line(sku_data, x='date', y='competitor_stock', title=f'Competitor Stock Over Time - {sku_select}')
fig2.add_scatter(x=sku_data[sku_data['low_stock_alert']]['date'],
                 y=sku_data[sku_data['low_stock_alert']]['competitor_stock'],
                 mode='markers', name='Low Stock Alert', marker=dict(color='red', size=10))
fig2.show()

```



```
import pandas as pd
```